

When moment summaries can certify attention

Boundary certificates, value-aware bounds, and GH200 evidence.

Sam Mausberg

Research manuscript, version 0.2.0. September 5, 2026.

Abstract

Moment summaries can avoid a token scan when their remaining uncertainty cannot change the consumer’s observation. We characterize this condition exactly for independent block intervals, construct exponential-attention enclosures from projected signed moments, and strengthen them with the dependence between block mass and value range. A Lean development proves the real-arithmetic chain from concrete finite block witnesses to an observation cell; all 64 theorems build and pass an axiom audit. An exact-rational study finds 61 of 149 initial coordinates certifiable with value coupling versus 57 with boxes, including a controlled example that defeats a global value-range baseline. GH200 experiments explain why reducing arithmetic alone is insufficient: parallel block reduction improves a favorable resident screening stage from 118.65 to 20.09 microseconds, but its complete host-decision path takes 83.11 microseconds against a 35.62 microsecond fused dense baseline. Actual post-RoPE traces from Qwen2.5-0.5B produce no whole-output passes in 2,016 rank-eight head instances; discarded scores dominate, while full-rank checks expose exponential remainder growth. The resulting contribution is a proved certificate interface and an empirical account of its information and execution bottlenecks.

Artifact and claim scope. The machine-checked result concerns real attention under explicit centering, radius, symmetry, and row-sum witnesses. The rational executable and its polygon optimizer are independently tested implementations; they are not extracted from Lean. The directed CUDA checker consumes outward imported rational boxes. The fast NumPy/CUDA summary path supplies numerical screens, whose floating-point soundness is not proved. Exact-real rounding also has a different contract from a particular fused GPU kernel. The [formalization map](#), [execution specification](#), and [recorded artifacts](#) give the precise boundaries.

Research questions

We ask what information a summary must retain to decide an observation, which term prevents that decision, and when avoided token reads repay construction and screening. These questions motivate three distinct comparisons: boxes against coupled metadata, analytic uncertainty against actual score variation, and an optimized resident kernel against its complete execution path. The negative cases determine the method’s useful scope as directly as the positive examples.

1. Definitions and the observation contract

Let the visible input be a finite, nonempty set partitioned into nonempty blocks $b = 1, \dots, B$. Write a positive normalized reduction as

$$y_j = \frac{N_j}{Z}, \quad Z = \sum_b z_b > 0, \quad N_j = \sum_b n_{bj}.$$

For a fixed block center ν_{bj} , define the centered numerator

$$m_{bj} = n_{bj} - \nu_{bj} z_b.$$

Suppose metadata establishes sound intervals

$$0 \leq l_b \leq z_b \leq u_b, \quad m_{bj}^- \leq m_{bj} \leq m_{bj}^+.$$

All intervals refer to the same input, mask, query, scaling, and version of the underlying data. Positivity of Z is a separate premise. In executable checks, $\sum_b l_b > 0$ is a sufficient way to discharge it.

An observation cell (a_j, b_j) is an open interval on which the consumer is constant. For finite numerical BF16 rounding, one may choose the midpoints to the preceding and following representable values. Strict inequalities avoid tie-breaking ambiguity. This formulation models numerical values, not the bitwise distinction between positive and negative zero.

Lemma 1 (constant consumer). If $y \in S$ and $C(x) = c$ for every $x \in S$, then $C(y) = c$.

Proof. Substitute y in the premise. For monotone rounding, an alternative closed-interval test is $Q(l) = Q(u)$ with $l \leq y \leq u$; monotonicity gives $Q(l) \leq Q(y) \leq Q(u)$. $\square$

Lean source: `constant_consumer`, `monotone_rounding`. A concrete IEEE-754 model and its monotonicity theorem are not supplied.

2. Exact residual certificates over block boxes

At a boundary c , define

$$\begin{aligned} \mathcal{L}_j(c) &= \sum_b \left[m_{bj}^- + \min\{(\nu_{bj} - c)l_b, (\nu_{bj} - c)u_b\} \right], \\ \mathcal{U}_j(c) &= \sum_b \left[m_{bj}^+ + \max\{(\nu_{bj} - c)l_b, (\nu_{bj} - c)u_b\} \right]. \end{aligned}$$

Theorem 2 (boundary certificate). Under the interval and positivity premises,

$$\mathcal{L}_j(c) \leq N_j - cZ \leq \mathcal{U}_j(c).$$

Consequently,

$$\boxed{\mathcal{L}_j(a_j) > 0 \quad \text{and} \quad \mathcal{U}_j(b_j) < 0 \quad \implies \quad a_j < y_j < b_j.}$$

Proof. The exact residual identity is

$$N_j - cZ = \sum_b \{m_{bj} + (\nu_{bj} - c)z_b\}.$$

For a fixed scalar α , the extrema of αz on $[l, u]$ are $\min(\alpha l, \alpha u)$ and $\max(\alpha l, \alpha u)$. Apply this identity with $\alpha = \nu_{bj} - c$, add the centered-numerator bounds, and sum. Dividing the first strict sign inequality by the positive Z gives $a_j < y_j$; the other gives $y_j < b_j$. $\square$

Lean source: `mul_enclosure`, `residual_identity`, `residual_enclosure`, `observation_cell`.

This test does not divide to evaluate the certificate. A numerical division may still be used to choose a candidate cell. A poor candidate cannot invalidate a sound certificate: it only makes certification less likely.

Theorem 3 (optimality for the stated metadata). Treat the block intervals as an independent Cartesian product, with no extra coupling constraints. For every fixed c , $\mathcal{L}_j(c)$ and $\mathcal{U}_j(c)$ are the attained minimum and maximum of the residual over that product.

Proof. For the minimum, choose $m_{bj} = m_{bj}^-$ and choose $z_b = l_b$ when $\nu_{bj} - c \geq 0$, otherwise $z_b = u_b$. These choices are simultaneous, admissible vertices of the product box and make every summand its minimum. The opposite choices and $m_{bj} = m_{bj}^+$ attain the maximum. Theorem 2 supplies the matching inequalities. $\square$

Lean source: `lowerChoice_mem`, `upperChoice_mem`, `lowerChoice_attains`, `upperChoice_attains`, `lower_attained`, `upper_attained`.

This is not a universal optimality claim about attention algorithms. Actual block masses and centered numerators can be correlated; the product box forgets that information. Richer metadata may give tighter bounds. The two extrema need not be attained by the same input. If the extremizers are interpreted as normalized outputs rather than residuals, require $\sum_b l_b > 0$ so that all vertices have positive denominator. Residual-box optimality itself does not need that additional condition.

Example. Take two blocks with centers 0, 2, masses in $[1, 2]$, and centered numerators identically zero. The true output range is $[2/3, 4/3]$. The boundary certificate proves membership in $(0.6, 1.4)$. A symmetric residual bound around 1, using midpoint masses, yields radius 0.5 and does not prove this cell. This establishes a strict improvement for this example, not domination of every possible earlier correlated certificate.




Figure 1. An analytic two-block example. The signs at 0.6 and 1.4 prove the required open cell. The output range is closed because its extrema are attained; the illustrative cell is not a BF16 rounding cell.

Corollary 3.1 (when the metadata is sufficient). If $\sum_b l_b > 0$, every output compatible with the independent block box lies in (a_j, b_j) if and only if $\mathcal{L}_j(a_j) > 0$ and $\mathcal{U}_j(b_j) < 0$.

Proof. Sufficiency is Theorem 2. If the lower sign test fails, Theorem 3 supplies a box vertex whose residual at a_j is nonpositive. Its denominator is positive, so its normalized output is at most a_j . The upper extremizer gives the corresponding counterexample at b_j . $\square$

Lean: `observation_cell_iff`. This converse explains a rejection precisely: the box metadata admits a conflicting output. It does not say that this conflicting vertex is realizable by the original tokens. Improving acceptance therefore requires tighter intervals, additional dependence information, or a different observation contract.

3. Monotone local refinement

Theorem 4 (refinement). Replace each mass interval by a nonempty subinterval and each centered-numerator interval by a nonempty subinterval. For every fixed boundary c , $\mathcal{L}_j(c)$ cannot decrease and $\mathcal{U}_j(c)$ cannot increase. If both the old and new enclosures contain the truth, their intersection does too.

Proof. A minimum over a smaller set cannot be smaller; a maximum cannot be larger. Equivalently, apply the endpoint multiplication bounds to both new endpoints, then sum. Intersection contains any point common to its two operands. $\square$

Lean source: `term_refinement`, `residual_refinement`, `interval_intersection`.

A scheduler may therefore refine selected blocks rather than discard all summary work when one test fails. The invariant is sound enclosure, not the particular choice of priority. The implementation's uncertainty priority is a heuristic, not an optimal scheduling theorem. Monotonicity concerns a fixed cell; changing the candidate cell changes the question being tested.

No unconditional termination with an open-cell certificate is asserted. A true output at a rounding midpoint may remain ambiguous forever under interval refinement. A complete executor needs a tie-aware path or a declared reference fallback, and a cap on refinement work.

3.1 Preserving mass and value dependence

For a block/channel, retain the exact centered value extrema $p_{bj} = \min_i(v_{ij} - \nu_{bj})$ and $q_{bj} = \max_i(v_{ij} - \nu_{bj})$. Positive weights imply

$$p_{bj}z_b \leq m_{bj} \leq q_{bj}z_b.$$

This follows by multiplying each centered-value inequality by its nonnegative weight and summing. Lean proves this step as `centered_value_mass_coupling`. It retains dependence that a separate mass/numerator box discards.

Define the nonempty compact polygon

$$P_{bj} = \{(z, m) : l_b \leq z \leq u_b, \quad m_{bj}^- \leq m \leq m_{bj}^+, \quad p_{bj}z \leq m \leq q_{bj}z\}.$$

Theorem 4.1 (coupled certificate). Replace each block summand in $\mathcal{L}_j(c)$ by $\min_{(z,m) \in P_{bj}}[m + (\nu_{bj} - c)z]$, and each summand in $\mathcal{U}_j(c)$ by the corresponding maximum. The resulting certificate is sound, never wider than the box certificate, and improves monotonically under sound box intersection with fixed value extrema.

Proof. The true mass and centered numerator belong to P_{bj} . The residual identity from Theorem 2 still holds. Each polygon is contained in its original rectangle, so its minimum cannot be smaller and its maximum cannot be larger. Shrinking the rectangle shrinks its intersection with the fixed cone. Summing and dividing by the positive total mass proves the claim. $\square$

The rational implementation evaluates this support exactly. A linear function on a nonempty compact polygon reaches its extrema at vertices. Candidate mass coordinates are l_b, u_b and $m_{bj}^-/s, m_{bj}^+/s$ for each nonzero slope $s \in \{p_{bj}, q_{bj}\}$. At a feasible candidate z , the centered endpoints are $\max(m_{bj}^-, p_{bj}z)$ and $\min(m_{bj}^+, q_{bj}z)$. These include all vertices: they are intersections of vertical, horizontal, or sloping boundary lines. The two sloping lines meet at $z = 0$, which, if feasible for nonnegative mass, is a mass endpoint. Coincident lines and degenerate segments need no additional candidates. Inconsistent metadata produces an empty candidate set and is rejected. This finite algorithm has a constant number of rational operations per block/channel; its bit cost depends on the fractions.

Joint optimization can outperform intersecting two independently computed scalar residual bounds. With $z \in [1, 4]$, $m \in [1/2, 3/2]$, $-z/2 \leq m \leq z/2$, and objective $m - z/4$, the box gives $[-1/2, 5/4]$ and the cone with the mass interval gives $[-3, 1]$. Their scalar intersection is $[-1/2, 1]$; joint optimization gives $[-1/2, 3/4]$. This is a metadata geometry example, not an assertion that every vertex is a token realization.




Figure 2. A second analytic example shows how value dependence excludes the box corners that prevent both signs. The required interval is illustrative.

The Lean library proves the positive-weight coupling and the box certificate. The polygon support algorithm is supported by the argument above and independent exact geometry tests; its executable correctness is not formalized.

3.2 A quantitative condition for refinement

At a fixed boundary c , the width of the box residual is

$$W_j(c) = \mathcal{U}_j(c) - \mathcal{L}_j(c) = \sum_b [m_{bj}^+ - m_{bj}^- + |\nu_{bj} - c|(u_b - l_b)].$$

Proposition 4.2 (margin versus uncertainty). For the true output inside (a_j, b_j) , it suffices that $W_j(a_j) < Z(y_j - a_j)$ and $W_j(b_j) < Z(b_j - y_j)$.

Proof. At the lower boundary the true residual is $Z(y_j - a_j) > 0$. Since it lies in $[\mathcal{L}_j, \mathcal{U}_j]$, the lower endpoint is at least the true residual minus W_j . The upper-boundary argument is symmetric. $\square$

This is an explanatory sufficient condition, not an executable rule requiring knowledge of the true output. It shows why near-midpoint coordinates are hard, why refining the block with the largest residual width is sensible, and why smaller numerical error alone does not establish a rounding decision. If sound refinement widths converge to zero and the true output has positive distance from both boundaries, a fixed cell eventually passes. The present bounded executor does not assume that this convergence occurs before fallback.

4. Projected signed-moment attention enclosures

4.1 Stored quantities

Attention is

$$y_j(q) = \frac{\sum_i e^{q^T k_i} v_{ij}}{\sum_i e^{q^T k_i}}.$$

The query already includes attention scaling. Keys are those actually attended to after positional transformations. Let μ_b, ν_b be exact block means, $\delta_i = k_i - \mu_b$, and $x_i = v_i - \nu_b$.

Select $r \leq d$ coordinates. The implementation uses a coordinate prefix; permutation can express another fixed subset. It does not silently substitute an approximate PCA basis. Split

$$q^T \delta_i = t_i + e_i, \quad t_i = u^T \delta_i^{(r)}, \quad |t_i| \leq \rho_b, \quad |e_i| \leq \epsilon_b.$$

Store coordinate radii $a_{bk} \geq \max_{i \in b} |\delta_{ik}|$ and value radii $R_{bj} \geq \max_{i \in b} |x_{ij}|$. Valid query bounds are

$$\rho_b = \sum_{k < r} |q_k| a_{bk}, \quad \epsilon_b = \sum_{k \geq r} |q_k| a_{bk}.$$

Store

$$\Sigma_b = \text{mean}(\delta^{(r)} \delta^{(r)T}), \quad C_b = \text{mean}(\delta^{(r)} x^T), \quad H_{bj} = \text{mean}(\delta^{(r)} \delta^{(r)T} x_j).$$

The full H tensor is temporary construction data, not query metadata. Retain its diagonal D_{bj} and a scalar witness

$$\eta_{bj} \geq \max_a \sum_k |(H_{bj} - D_{bj})_{ak}|.$$

Here D_{bj} denotes the diagonal matrix or its stored diagonal as appropriate. The choice $D = 0$ is also valid. Removing the diagonal weakly reduces the stored absolute row-sum witness, but also shifts the approximation center. The resulting intervals need not be nested, so acceptance need not improve for every cell. For example, take $H_{11} = 9$, $H_{23} = H_{32} = 10$, all other entries zero, and $u = (1, 0, 0)$. Both row-sum witnesses equal 10. The quadratic enclosure changes from $[-10, 10]$ to $9 + [-10, 10]$. The implementation exposes both representations. For $r = 0$, define the empty row-sum witness to be zero.

Lemma 5 (quadratic witness). If F is symmetric and $\eta \geq \max_a \sum_k |F_{ak}|$, then

$$|u^T F u| \leq \eta \|u\|_2^2.$$

Proof. Use $|u_a u_k| \leq (u_a^2 + u_k^2)/2$. Symmetry makes the row and column contributions equal, hence

$$|u^T F u| \leq \frac{1}{2} \sum_{a,k} |F_{ak}| (u_a^2 + u_k^2) = \sum_a u_a^2 \sum_k |F_{ak}| \leq \eta \sum_a u_a^2. \quad \square$$

This proof avoids numerical eigenvalue computations. Lean proves it as `quadratic_rowsum_bound`; `signed_moment_omission_bound` instantiates it on the actual finite signed tensor and its retained symmetric part.

4.2 A sharper Taylor majorant

Define

$$\kappa(\rho) = \begin{cases} (e^\rho - 1 - \rho - \rho^2/2)/\rho^2, & \rho > 0, \\ 0, & \rho = 0. \end{cases}$$

Lemma 6 (second-order exponential remainder). For $\rho \geq 0$ and $|t| \leq \rho$,

$$|e^t - 1 - t - t^2/2| \leq \kappa(\rho) t^2, \quad \kappa(\rho) \leq e^\rho \rho/6.$$

The latter inequality is strict for $\rho > 0$. For positive radius, the first coefficient is the smallest uniform quadratic majorant on this interval: $t = \rho$ attains equality. At $\rho = 0$ we choose the nonnegative coefficient zero; minimality over arbitrary real coefficients is not asserted.

Proof. At $\rho = 0$, $t = 0$. Otherwise absolute convergence of the exponential series gives

$$\left| \sum_{n \geq 3} \frac{t^n}{n!} \right| \leq t^2 \sum_{n \geq 3} \frac{\rho^{n-2}}{n!} = \kappa(\rho) t^2.$$

Writing $n = m + 3$ and using $(m + 3)! \geq 6m!$ yields the second inequality by termwise comparison with $(\rho/6)e^\rho$. For $m \geq 1$ the factorial inequality is strict and $\rho^{m+1} > 0$. $\square$

The exact-real inequality is sharper than the earlier Lagrange bound. Near zero, the implementation evaluates a series rather than subtracting nearly equal floating values. That implementation is numerical unless an outward remainder is included. The rational path uses an upper enclosure of e^ρ .

Lean derives the analytic bound from the convergent `Real.exp` series in `exp_remainder_sharp`. `exp_remainder_endpoint`, `exp_coefficient_optimal`, and `exp_coefficient_lt_lagrange` establish endpoint attainment, optimality, and strict improvement. The conditional lifting lemma is then used with its remainder premise discharged by the analytic theorem.




Figure 3. The sharper coefficient improves constants substantially at large radius, but still grows exponentially. Projection has a separate exponential cost. The BF16 line is a unit-scale guide rather than a sufficient certificate.

4.3 The enclosure theorem

Fix any common finite shift s and define the actual block quantities by

$$z_b(s) = \sum_{i \in b} e^{q^T k_i - s}, \quad m_{bj}(s) = \sum_{i \in b} e^{q^T k_i - s} x_{ij}.$$

The common factor e^{-s} cancels from the normalized output. All subsequent endpoints and residuals refer to these shifted quantities. Set $a_b = q^T \mu_b$, $w_b = n_b e^{a_b - s}$, and

$$\sigma_b^2 = u^T \Sigma_b u, \quad A_b = 1 + \sigma_b^2/2, \quad \tau_b = \kappa(\rho_b) \sigma_b^2.$$

For the projected scores, define

$$\begin{aligned}\tilde{l}_b &= w_b \max\{1, A_b - \tau_b\}, & \tilde{u}_b &= w_b(A_b + \tau_b), \\ \widehat{m}_{bj} &= w_b \left[(C_b^T u)_j + \frac{1}{2} u^T D_{bj} u \right], & \beta_{bj} &= w_b \left[\frac{1}{2} \eta_{bj} \|u\|_2^2 + \tau_b R_{bj} \right].\end{aligned}$$

Theorem 7 (full-score block enclosures). The following are sound enclosures for the actual, unprojected attention block:

$$\boxed{l_b = e^{-\epsilon_b} \tilde{l}_b, \quad u_b = e^{\epsilon_b} \tilde{u}_b,}$$

$$\boxed{m_{bj}^\pm = \widehat{m}_{bj} \pm \left[\beta_{bj} + (e^{\epsilon_b} - 1) \tilde{u}_b R_{bj} \right].}$$

Proof. Centering gives $\text{mean}(t_i) = 0$, $\text{mean}(t_i^2) = \sigma_b^2$ and $\text{mean}(x_i) = 0$. By Lemma 6, the mean absolute remainder is at most τ_b . Consequently the projected mass is in $w_b[A_b - \tau_b, A_b + \tau_b]$. Jensen's inequality also makes it at least w_b , proving the asserted projected mass interval.

The exact projected centered numerator is

$$w_b \left[(C_b^T u)_j + \frac{1}{2} u^T H_{bj} u + \text{mean}(r_i x_{ij}) \right],$$

where $r_i = e^{t_i} - 1 - t_i - t_i^2/2$. Lemma 5 bounds the omitted quadratic form; $|\text{mean}(r_i x_{ij})| \leq \tau_b R_{bj}$ bounds the other omission. This proves the projected centered interval.

Restoring discarded coordinates multiplies each positive projected weight by $e^{e_i} \in [e^{-\epsilon_b}, e^{\epsilon_b}]$, giving the mass enclosure. Moreover $|e^{e_i} - 1| \leq e^{\epsilon_b} - 1$, so the absolute change in the centered numerator is at most

$$(e^{\epsilon_b} - 1) R_{bj} \sum_{i \in b} e^{a_b - s + t_i} \leq (e^{\epsilon_b} - 1) \tilde{u}_b R_{bj}.$$

Add this perturbation to the projected interval. $\square$

Lean proves this chain in `full_score_enclosure`, then includes the common score offset in `moment_block_enclosure`. `full_attention_observation` composes these block results with the residual signs into the original finite per-token weighted quotient. Its witness contains centering, radius, symmetry, and row-sum conditions; it does not assume the Taylor remainder or the desired output enclosure. Lean uses unnormalized finite-sum tensors: multiply the paper's H, D, η by the block cardinality to obtain that representation. The score decomposition and identification of the intended input partition must still match the caller's data; the formal theorem is not a Python summary validator.

The cases $r = 0$ and $r = d$ are included. At $r = 0$, the retained centered moment vanishes and all score variation is bounded as discarded variation. At $r = d$, $\epsilon_b = 0$ and no projection inflation is paid. Neither projection nor a fixed rank is assumed to be useful on arbitrary trained-model activations.

5. Rational checking and outward export

`cmk/rational.py` evaluates the summaries with exact Python fractions and bounded exponential series. For $0 \leq x \leq 1/2$, after summing through degree n , the first omitted term is $a_{n+1} = x^{n+1}/(n+1)!$. Every subsequent term ratio is at most $q = x/(n+2) < 1$, so the tail is at most $a_{n+1}/(1-q)$. The partial sum and partial sum plus this tail enclose e^x . Positive interval squaring

handles range reduction; reciprocal intervals handle negative inputs. All arithmetic in this procedure is rational. The implementation restricts its input domain rather than returning an unbounded or silently approximate result.

A direct rational exponential-sum oracle does not use moment formulas. Block refinement intersects its interval result with the existing enclosure. Exact fractions are expensive; these routines are correctness artifacts, not the proposed fast GPU implementation.

`cmk/export.py` converts each rational endpoint outward to binary64 and compares the converted float’s exact rational value against the original. This also exports intervals for the centers, whose rational means may not be exactly representable. The host checker conservatively expands basic operations by one ULP. The CUDA draft uses directed-rounding arithmetic. Their validity still requires the stated floating-point semantics and faithful compilation. Both programs have now been compiled and executed on this GH200, including the imported CUDA path under memory checking. Neither is extracted from Lean.

It is invalid to pass the ordinary NumPy/CUDA moment estimates to this checker and assume that directed rounding at the final sum repairs missing uncertainty in the input metadata. Sound inputs are a necessary premise.

6. Generalization to consumers and smooth gates

Theorems 2 through 4 require only positive normalized weights, not exponentials, attention heads, or a particular network. A new weight function must supply its own sound block envelopes. Signed, unnormalized contractions can instead sum intervals directly.

For a smooth gate ϕ and scalar deviations α, γ , set $p = \phi(t + \alpha)$, $p_0 = \phi(t)$, and $p_1 = \phi'(t)$. Then

$$p(s + \gamma) - (p_0 s + p_1 \alpha s + p_0 \gamma) = p_1 \alpha \gamma + (p - p_0 - p_1 \alpha)(s + \gamma).$$

If $|\alpha| \leq \rho_g$, $|\gamma| \leq \rho_u$, and a Taylor witness supplies $|p - p_0 - p_1 \alpha| \leq M \rho_g^2 / 2$, the triangle inequality gives

$$|\text{error}| \leq |p_1| \rho_g \rho_u + \frac{1}{2} M \rho_g^2 (|s| + \rho_u).$$

Lean source: `gated_residual_identity`, `gated_remainder_bound`. The analytic curvature bound for a concrete gate is an additional obligation.

For SwiGLU, this identity applies to each neuron before the down-projection. Summing absolute downstream weights times these nonnegative local error radii gives a componentwise output bound. Contracted coefficients can avoid a hidden intermediate, but may cost more storage than the original weights. No useful general feed-forward compression or dense-GEMM acceleration is established here.

For an argmax consumer, $l_w > u_j$ for every $j \neq w$ proves the winner w . For a deterministic stateful network, equality of each complete state transition implies equality of repeated execution. These statements do not turn local approximation bounds into whole-model equality.

Lean source: `strict_argmax`, `transition_composition`.

7. Representation, execution, and the cost of reuse

The base box representation stores

$$B(2d + r^2 + 2rh + 3h + 1)$$

floating scalars: key centers and radii, projected covariance, key/value cross moments, signed diagonals, value centers and radii, row-sum witnesses, and counts. The Python coupling extension adds

$2Bh$ scalars for value extrema. These counts exclude original K/V, index maps, source identities, scratch, and allocations. The benchmark records the CPU summary size and the actual uploaded GPU arrays separately. Keeping original K/V available for fallback remains necessary.

A shared block-scalar pass computes centroid scores, variance, score radii, and mass terms once per query/block. A channel pass contracts value moments; a reduction performs the boundary tests. This gives query arithmetic $O(B(d + r^2 + rh + h))$, compared with $\Theta(N(d + h))$ for a direct token scan. These are arithmetic/word-operation counts. They omit rational bit complexity and do not imply a timing ratio. The CPU builder scans full K/V and constructs the temporary signed tensor, costing $O(N(d + h + r^2h))$ with the present method. Supplied synthetic blocks require no clustering; clustering an actual workload would add another measured setup stage.

The first CUDA evaluator repeated scalar work across channels. We implemented a shared scalar pass and then measured a different bottleneck: serial reduction over blocks. A third variant assigns a CTA to an output coordinate and reduces blocks in parallel. This changes the summation order of the numerical path. The directed imported-box checker remains separate. Section 10 measures all three variants under matched inputs, including configurations where extra parallelism hurts.

Proposition 8 (reuse threshold in a fixed-cost model). Let construction cost S , dense query cost D , full-output pass probability p , and expected screen/decision/return overhead C . If rejected queries execute the full dense path, expected total time for R reuses is

$$T(R) = S + R[C + (1 - p)D].$$

It beats RD precisely when $pD > C$ and $R > S/(pD - C)$.

Proof. Subtract RD and rearrange. $\square$

This model assumes fixed dense cost and accounts for decision and return work in C ; a heterogeneous scheduler requires its actual conditional costs. It explains why high pass rate and small summaries can coexist with a slow system. In our measured complete implementation the steady query path is already slower than dense, so no finite amortization count rescues it. We report the measured path directly instead of inferring its latency from p .

The executable protocol is bounded: validate the supported input and identity, evaluate metadata, choose a candidate, check strict signs, then refine or invoke the declared fallback. Exact-rational refinement intersects sound intervals. Numerical CUDA correction scans are experimental estimates. Source identities bind Python refinement to the exact ordered visible data, query and shift. They do not implement concurrent cache epochs, RoPE updates, or a serving memory manager. The [GH200 specification](#) defines those boundaries.

8. Correctness evidence and the value of coupling

8.1 Universal claims and executable checks

The Lean 4.24.0 build and complete axiom audit cover 64 local declarations, including the actual exponential remainder, finite signed tensor contractions, full-score block enclosure, and composed real-attention quotient. Only `propext`, `Classical.choice`, and `Quot.sound` occur. Source hashes and per-theorem dependencies are retained in [lean-verification.json](#). The polygon optimizer, rational exponential routine, BF16 cell constructor, state identity, and instruction-level CUDA behavior are implementation obligations beyond that theorem chain.

The 46-test CPU suite includes the original randomized enclosure tests and new independent checks: 500 exact polygon comparisons against a generic half-plane intersection implementation, 40 randomized direct 100-digit attention comparisons across rank and refinement, and 40 numerical/exact geometry comparisons. It also covers invalid dimensions, nonintegral indices, negative bounds, unsupported domains, stale source/query/shift, and midpoint behavior. These tests target distinct mathematical and input contracts.

The audit exposed an actual arithmetic defect: constructing `Interval(3, 3)` with integers allowed reciprocal evaluation through Python’s floating-point `1 / int`. Thus `Interval(1,1) / Interval(3,3)` could have an upper endpoint below exact $1/3$. Endpoints and scale factors now normalize to exact fractions, and the regression requires both division endpoints to equal $1/3$ exactly. This example motivates auditing the executable arithmetic independently of a correct real-number theorem.

The imported CUDA checker processes 160 rationally generated rows. It retains all 74 rational-certified rows, produces no certificate on a rationally refused row in that fixture, and rejects 13 explicit invalid-domain or strict-boundary controls. Numerical CUDA fixtures compare original/shared/parallel evaluation and dense block correction against CPU results. Memory, race, and synchronization checks report no errors on the exercised numerical paths; imported CUDA memory checking also reports no errors. This is finite implementation evidence under the documented compiler and rounding semantics.

8.2 What additional value metadata actually buys

The exact-rational coupling study uses seed 962605, four retained ranks and four key spreads, narrow and broad value channels, exact midpoint channels, and named geometry and attention witnesses. All 53 cases and each sequential refinement stage are saved, for 149 initial coordinates.

Initial result	Count
Box certificate accepts	57 / 149
Coupled certificate accepts	61 / 149
Global value-hull baseline accepts	50 / 149
Strictly narrower coupled residual	30 / 149
Unchanged width	119 / 149
Added coupling accepts beyond the global hull	1
Observed false certificate or independent coverage failure	0

The global hull matters as an explanation. If every value already lies in one rounding cell, positive weights alone certify the result. Three of the four extra coupling accepts have that explanation. The remaining prescribed witness has keys $-2, 2$ and values $1 - 1/2560, 1 + 1/2560$ in one rank-zero block, plus a singleton key -16 with value 3, at query $q = 1$. The global hull spans the outlier and fails. Its very small block mass allows the coupled block test to certify BF16 value 1. An independent direct rational exponential sum verifies this certificate. The mechanism is local value range combined with mass, not a claim that four selected examples estimate model-wide coverage.




Figure 4. The full initial ablation retains unchanged widths and refused coordinates. The global-hull baseline explains most added accepts. Raw inputs, intervals, refinement stages, setup, and actual rational fallback costs are in [coupling.json](#).

Independent high-precision comparisons use 110 decimal digits and an explicit scaled 10^{-95} tolerance for numerical comparison with exact zero-width intervals. The first zero-tolerance development attempt failed on cancellation in a constant channel; that methodological failure is documented. Certificate and rational-oracle comparisons use exact fractions without this tolerance. Stored decimal endpoints are displays rather than reusable certificate data. Broad channels and exact midpoint abstentions remain in the artifact.

9. Why the bound succeeds or fails

9.1 Local uncertainty and selective refinement

A numerical CPU experiment uses $N = 8192, d = 16, h = 8, B = 16, Q = 24$ and one BLAS thread. The generator supplies favorable four-coordinate structure. With tight blocks, the metadata-only screen passes all 24 queries. Introducing one broad block causes every initial query to fail; correcting that block evaluates the attention weights and values of exactly $1/16$ of the tokens. Making all blocks broad requires correction of every token. The correction-token counter excludes the current Python executor’s full-source input validation and fingerprint passes, which read K/V even when no correction is needed. Its actual CPU timing includes those passes, so this counter is not a total memory-traffic measurement. The experiment identifies uncertainty localization rather than GPU speedup. The complete Python paths are slower than dense NumPy, with construction and actual fallback costs retained.




Figure 5. One uncertain block permits local correction; diffuse uncertainty removes that advantage. Every query’s correction-token fraction is shown. Full-source validation reads are additional and included in the reported CPU times.

The GPU synthetic controls keep $d = h = 64$ and use BF16 input values represented exactly in binary64 for the summary path. Tight clusters can pass fully. Aggregating seeded inputs reused across ranks, moderate cases pass 10,197 of 19,008 coordinates across configurations but none of 297 complete outputs. In the centered-value control, 1,986 of 2,048 coordinates pass while only 6 of 32 outputs do. An output is a conjunction of coordinate decisions; high marginal coordinate coverage is therefore a weak predictor of complete-output usefulness. Independence between coordinates is neither assumed nor needed for this observation. Exact-midpoint controls reject all strict cells by design.




Figure 6. These rank-four, $N = 8192$, $Q = 32$ controls change the source of uncertainty and the output’s cell margins. One failed coordinate triggers the measured full-batch fallback.

9.2 Trained post-RoPE activations

We capture actual Q/K/V arguments at the attention call of the public Qwen2.5-0.5B model [13], pinned to revision 060db6499f32faf8b98477b0a26969ef7d8b9987, using Transformers 4.51.3 and BF16 weights. The arrays are captured after RoPE. Three deterministic original probe texts cover descriptive prose, code, and arithmetic at prefix lengths 1,024 and 4,096. The final prefill query sees the full supplied prefix, with contiguous GQA groups of seven query heads per KV head. Scaling by $1/8$ is exact for the captured values. This setup is an activation diagnostic, not a language-model accuracy benchmark or a sample of production traffic.

All 24 layers and 14 query heads are checked at retained rank eight: six prefixes times 336 heads gives 2,016 head instances and 129,024 coordinates. We additionally check full rank in layers 0, 12 and 23, giving 252 head instances. Both sets produce zero numerical coordinate or whole-output passes. Retained coordinates are a prefix of the actual transformed basis. No learned basis, PCA fitting, or clustering was performed, so this result does not rule out other representations.

Radius diagnostic	Rank 8	Rank 64, representative layers
Median maximum-block retained bound ρ	1.32	29.91
Median maximum-block discarded bound ϵ	34.24	0
Median actual maximum projected score deviation	0.85	8.91
Median minimum-channel BF16 boundary margin	$8.95 \cdot 10^{-7}$	$6.66 \cdot 10^{-7}$




Figure 7. Low rank leaves substantial discarded-score variation. At full rank, coordinate bounds are looser than scan-derived radii, and the actual score variation is still substantial. Computing the diagnostic radii reads tokens; it is not free query metadata.

The failure has two levels. At rank eight, $e^\epsilon - 1$ amplifies the centered-value perturbation far beyond the cell margin. At full rank that term vanishes, but $\kappa(\rho)$ becomes large because the retained score radius is broad. The median maximum absolute candidate error is about 0.82 at rank eight and 0.50 in the full-rank subset. Even a tighter scalar bound would have to resolve every coordinate within much narrower cells. Figure 8 normalizes the omitted quadratic, Taylor, and projection terms by the actual shifted mass to make their scale interpretable in output units.




Figure 8. Box plots retain the distribution of the maximum-channel centered numerator terms divided by true shifted mass, alongside minimum-channel cell margins. The terms are explanatory diagnostics, not an exact additive decomposition of output error; mass uncertainty also contributes. Outliers are retained in the JSON even when excluded from the plot's whiskers.

A smaller rank reduces the stored rank-dependent arrays; useful certification also requires controlled discarded-score uncertainty. Increasing rank can replace projection inflation with Taylor inflation. This explains the unfavorable trace result more precisely than a single zero-coverage number. A useful next representation must address actual within-block score variation, the looseness of query-independent radius bounds, and value dependence together.

10. What the GH200 measurements establish

10.1 Matched execution and retained costs

The machine has one NVIDIA GH200 GPU, compute capability 9.0, 97,871 MiB reported GPU memory, driver 580.105.08, CUDA 12.8, and PyTorch 2.7.0. The final sweep contains 57 synthetic configurations and nine actual prose trace/rank configurations. It varies $N \in \{1024, 8192, 32768\}$, rank, key spread, and $Q \in \{1, 32\}$. The latter is a noncausal batch of queries sharing K/V, not 32 independent serving requests. Actual trace cases use seven query heads sharing one KV head at the final prefix position.

The baseline explicitly forces PyTorch’s fused FlashAttention SDPA backend, with the same BF16 input values, scale, visible keys, and no dropout. A separate binary64 softmax/matmul path checks numerical outputs. Resident query timings exclude setup by definition; setup, uploads, casts, allocations, workspace, and graph capture are recorded separately. Complete eager-path timings include screening, device reduction of decisions, CPU synchronization, and either an output cast or a full-batch dense fallback. Seven samples of 20 invocations supply raw CUDA-event and host-wall measurements. Plotted quantiles describe these timing samples, not a confidence interval for other workloads.




Figure 9. Wall-clock medians with 10th-90th sample quantiles for tight rank-four inputs. The complete path includes its actual decision and fallback behavior.

10.2 A bottleneck explanation, with an adverse case

Tight rank-four case	Original graph	Shared-scalar graph	Parallel-reduction graph	Dense graph
$N = 1024, Q = 32$	20.47 μs	19.87 μs	20.84 μs	9.80 μs
$N = 8192, Q = 1$	35.69 μs	35.60 μs	18.72 μs	14.55 μs
$N = 32768, Q = 1$	118.65 μs	118.14 μs	20.09 μs	31.30 μs
$N = 32768, Q = 32$	143.28 μs	147.52 μs	64.89 μs	34.21 μs




Figure 10. Sharing repeated scalar arithmetic barely changes the dominant serial block reduction. Exposing that block work across CTAs changes the long-context single-query case. Additional synchronization hurts at small N .

The original reduction runs one thread per output coordinate, each looping serially over blocks. A single query with 64 output coordinates exposes very little parallelism. A CTA per coordinate makes block parallelism available; for the favorable $N = 32768$, $Q = 1$ case, the isolated reduction phase falls from 99.38 to 6.65 microseconds. In contrast, 32 queries already expose more work, and the additional per-coordinate CTAs and synchronization can cost more than they save for small block counts. This is an architectural explanation tested by an implementation ablation, rather than a speedup predicted from arithmetic.




Figure 11. Isolated phase medians diagnose the bottleneck; their sum is not used as a substitute for measured complete latency. CPU construction, upload, and workspace allocation remain explicit.

10.3 Why the favorable kernel result is insufficient

For $N = 32768$, $Q = 1$, $r = 4$, the resident parallel screen takes 20.09 microseconds under fixed-input graph replay versus dense attention's 31.30 microseconds. The actual complete eager path instead takes 83.11 microseconds against 35.62 microseconds, with about 51.50 milliseconds of reusable summary and workspace setup. Its CPU moment-and-extrema array payload is 1,000,448 bytes, the uploaded box summary 869,376 bytes, and original BF16 K/V 8,388,608 bytes. This excludes 262,144 bytes of group indices and Python object/hash overhead. Scratch and original K/V retention are reported separately. The resident kernel advantage does not establish that the complete path saves time or memory in a serving system.




Figure 12. The complete query path is slower than dense in every measured configuration, so no finite reuse count repays setup under those measured costs. Fixed-input graph replay identifies remaining device work and omits the adaptive host decision; it is a separate experiment.

Across 66 final configurations, numerical accepted coordinates have no observed BF16 mismatch against the binary64 attention check, but 1,691 mismatches against the fused BF16 baseline. These aggregate counts include repeated configurations and are not independent accuracy trials. They illustrate the distinct targets: a real-attention rounding cell does not specify a fused kernel’s intermediate precision or accumulation. The binary64 comparison itself remains numerical, not an exact-real proof of each GPU decision.

The first trained-trace parity attempt also failed an absolute residual tolerance because broad bounds have extremely large magnitudes. The repaired comparison uses a 10^{-12} scaled tolerance and still requires identical screen decisions plus independent output checks. Maximum observed scaled variant difference is $1.33 \cdot 10^{-15}$. The initial nine failures and the original scalar-only sweep are retained. This tolerance is for implementation parity diagnostics; it is not added to a sound interval claim.

10.4 Reproduction at the committed implementation

The figures and tables retain the first complete sweep. A second run on the same GH200 measures commit d2c3668, including the final host ABI guards and formatted sources. Source, library, and trace hashes are recorded in [provenance.json](#); the original measurement is preserved alongside [reproduction.json](#). All 66 cases pass again, with unchanged coverage and mismatch counts. Across cases, parallel-screen graph medians differ by between -1.08% and $+1.42\%$, with a median ratio of 1.002. Complete-path and dense host-wall medians are noisier: their median ratios are 1.075 and 1.112, respectively. Every complete screened path remains slower than dense, and no finite setup break-even appears. This checks repeatability on this machine; it does not estimate variation across devices, model families, or serving systems.

11. Relation to primary literature

The relevant mathematical ingredients have established histories. Adaptive predicates decide a sign with a cheap filter and increase effort only when uncertainty remains [5]. Taylor models retain polynomial structure with a bounded remainder [9]. Our box extrema are elementary support functions of rectangles, and the coupled optimizer is elementary polyhedral optimization. The paper does not claim those principles as new.

In attention, Multipole Attention already clusters keys, performs selected exact work, and retains approximate contributions elsewhere [2]. TaylorShift rearranges polynomial attention, and symmetry-aware Taylor features exploit polynomial tensor structure [3,4]. COBS uses compressed covariance statistics to estimate block mass; SPLA combines second-order selection with a residual

linear path [11,12]. LOCKS also explicitly analyzes the limits of moment-only page summaries for broad or peaky scores [10]. The score-spread diagnosis here is consistent with that concern, while this paper measures its effect on strict output-cell certification. We have not reproduced those systems or claimed a performance advantage over them.

WitCert and attention-memory observability contracts already connect runtime witnesses, formal scope, and local fallback [6,7]. The distinction studied here is the residual interface for a specified consumer, the projected signed value moments supplying its inputs, and the stronger block mass/value coupling. The present 64-theorem development and GH200 experiments make that combination checkable and expose where it fails. They do not establish publication priority or a universally useful sublinear attention algorithm. The dated [primary-literature comparison](#) records specific overlap.

FlashAttention reorganizes dense computation around GPU memory traffic [1], and subsequent Hopper work exploits asynchronous execution [14]. Dense decode also admits sequence-parallel strategies [15]. Our forced PyTorch fused backend is a concrete comparison point; it is not an assertion that every newer or decode-specialized implementation has been exhausted. Consequently, the positive device-stage result has deliberately narrower scope than a claim of state-of-the-art attention performance.

12. Consequences and next research decisions

The measurements support a specific research direction. Better summaries must preserve information that controls the consumer decision, not only approximate attention mass. Value-range coupling demonstrates one such improvement. The trained traces separately expose large score uncertainty; coupled certification on those traces remains unmeasured. Block partitioning, score-radius witnesses, higher-order residuals, and consumer-specific margins are therefore separate design variables to test. Any query-dependent tightening that scans keys must account for that scan.

The GPU study identifies a separate implementation requirement. A resident parallel reduction can make the screening stage competitive in one regime, but removing the host decision and bounding fallback work are prerequisites for turning that result into a useful path. A sound GPU summary builder remains another prerequisite for claiming a rigorous accelerated executor. The present artifacts establish the real-arithmetic certificate, test a rational implementation, and measure numerical GPU behavior with those distinctions preserved.

References

- [1] T. Dao et al. FlashAttention: Fast and Memory-Efficient Exact Attention with IO-Awareness. 2022. [arXiv:2205.14135](#)
- [2] C. Hooper et al. Multipole Attention for Efficient Long Context Reasoning. 2025. [arXiv:2506.13059](#)
- [3] T. C. Nauen, S. Palacio, and A. Dengel. TaylorShift: Shifting the Complexity of Self-Attention from Squared to Linear (and Back) using Taylor-Softmax. 2024. [arXiv:2403.02920](#)
- [4] F. A. Heinsen and L. Kozachkov. Self-Attention at Constant Cost per Token via Symmetry-Aware Taylor Approximation. 2026. [arXiv:2602.00294](#)
- [5] J. R. Shewchuk. Adaptive Precision Floating-Point Arithmetic and Fast Robust Geometric Predicates. Discrete & Computational Geometry 18, 305-363, 1997. [Author's page](#)
- [6] F. Wei et al. WitCert: Sound Runtime Risk Observability and Gating for KV-Cache Quantization. 2026. [arXiv:2607.28699](#)
- [7] F. Wei et al. Runtime Observability for Heterogeneous Attention Memory. 2026. [arXiv:2608.05863](#)
- [8] Y. Kang, G. Tran, and H. De Sterck. Fast Multipole Attention: A Scalable Multilevel Attention Mechanism for Text and Images. 2023. [arXiv:2310.11960](#)

- [9] M. Berz and G. Hoffstätter. Computation and Application of Taylor Polynomials with Interval Remainder Bounds. *Reliable Computing* 4, 83-97, 1998. [doi:10.1023/A:1009958918582](https://doi.org/10.1023/A:1009958918582)
- [10] J. Hwang. Locks: Page-Local Compact Key Summaries for Efficient Long-Context Decoding. 2026. [arXiv:2607.24555](https://arxiv.org/abs/2607.24555)
- [11] A. Tian et al. COBS: Cumulant Order Block Sparse Attention. 2026. [arXiv:2607.09052](https://arxiv.org/abs/2607.09052)
- [12] B. Wang et al. SPLA: Block Sparse Plus Linear Attention for Long Context Modeling. 2026. [arXiv:2601.22379](https://arxiv.org/abs/2601.22379)
- [13] Qwen Team. Qwen2.5-0.5B model card and public weights. [Pinned checkpoint](#).
- [14] J. Shah et al. FlashAttention-3: Fast and Accurate Attention with Asynchrony and Low-precision. 2024. [arXiv:2407.08608](https://arxiv.org/abs/2407.08608)
- [15] R. Sanovar et al. Lean Attention: Hardware-Aware Scalable Attention Mechanism for the Decode-Phase of Transformers. 2024. [arXiv:2405.10480](https://arxiv.org/abs/2405.10480)